

Implementation details

rebalance()

- Use MPI_Allgather to let each process know the number of elements in each of the processes, store in sub_counts.
- Figure out the global number of elements and what is the rebalanced datasize *rCount in each process
- If a process currently has less load than *rCount, receive data from other processes until it reaches *rCount
- If a process has more load than *rCount, it analyzes the sub_counts array and figure out how many data it needs to send to which processes. All processes follow the algorithm as follows:
 - Starting from rank 0, scan the sub_counts array. Consider rank i:
 - If a processes j needs data, send $\min(\text{sub_counts}[i] - *r\text{Count}, *r\text{Count} - \text{sub_counts}[j])$ to it. Add the amount sent over to sub_counts[j] and remove from sub_counts[i]
 - Rank i continue scanning the array and send data until $\text{sub_counts}[i] == *r\text{Count}$

findSplitters() // assume data is already sorted

- Send local data size, local max to rank 0 using MPI_Reduce
- For rank 0 do the following:
 - Get an initial set of $\text{PROBE_MULTIPLE} * (\text{numSplitters} - 1)$ probes by using $\text{PROBE_MULTIPLE} * (\text{numSplitters} - 1)$ local values evenly spaced
 - Broadcast the number of probes to all processes
 - Broadcast the probes to all processes
 - Calculate ideal bin size = $\text{global_N} / \text{numSplitters}$, from the binsize we can derive the ideal splitter locations.
 - Count the number of elements in each bin separated by the probes
 - Get the global number of elements in each bin using MPI_Reduce
 - If any of the probes fall within 1% of the ideal splitter location, mark it as achieved.
 - For unachieved splitters, count how many splitters there are in each of the intervals separated by the probes. For an interval having m splitters, generate $\text{PROBE_MULTIPLE} * m$ new probes evenly spaced between the interval
 - Broadcast the number of probes and the probes to all processes
 - Iterate until all splitters are located.
- For all other ranks do the following:
 - Receive broadcasted number of probes from rank 0.
 - If number of probes > 0:
 - Receive broadcasted probes from rank 0
 - Calculate local probe ranks, send to rank 0 using MPI_Reduce
 - iterate
 - If number of probes == 0, it means all splitters has converged:
 - Receive broadcasted splitters array from rank 0

- Receive broadcasted global bin size from rank 0
- End function

moveData()

- All to all collective (MPI_Alltoall(), MPI_Alltoallv())

I also experimented with:

- One to all collective (MPI_Gather())
- Point to point messaging (MPI_Isend, MPI_Recv)

But they are all slower than all-to-all, possibly because collective communications are highly optimized in the MPI implementation.

Benchmarking

Target performance:

Benchmark results


```
Duration_Rebalance 0.225306 s
Duration_Splitters 0.00103572 s
Duration_Move 1.64638 s
Duration_Total 1.87272 s
Duration 1.87272 s
```

My performance running default batch.slurm

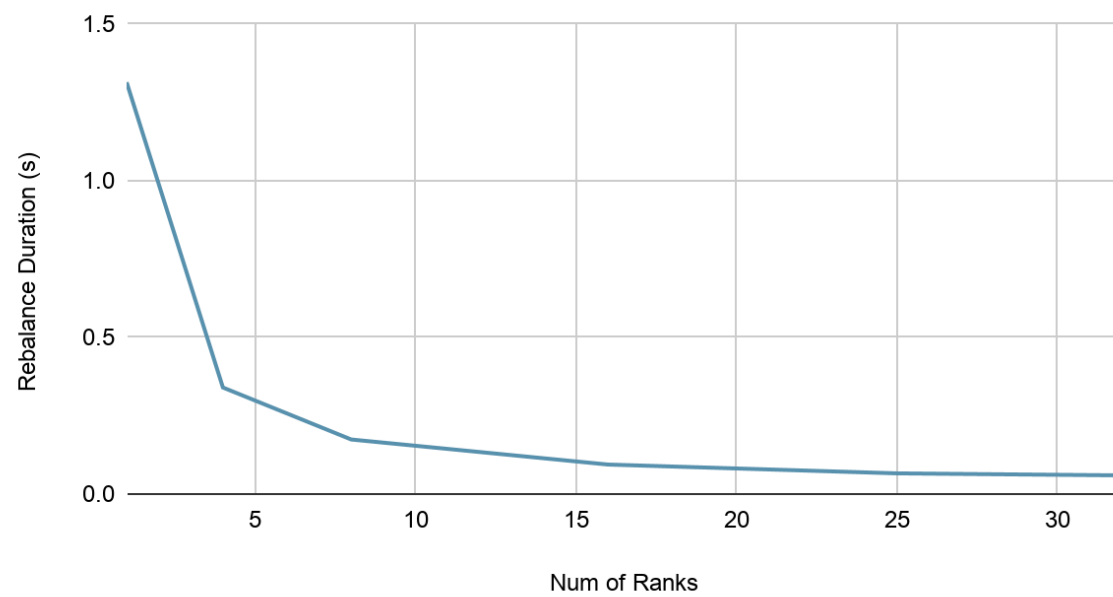
```
Benchmark results
Duration_Rebalance 0.0801739 s
Duration_Splitters 0.00722673 s
Duration_Move 1.39056 s
Duration_Total 1.47796 s
Duration 1.47796 s
```

Keeping the global data set at 500M and vary the number of ranks

Total Duration vs. Num of Ranks

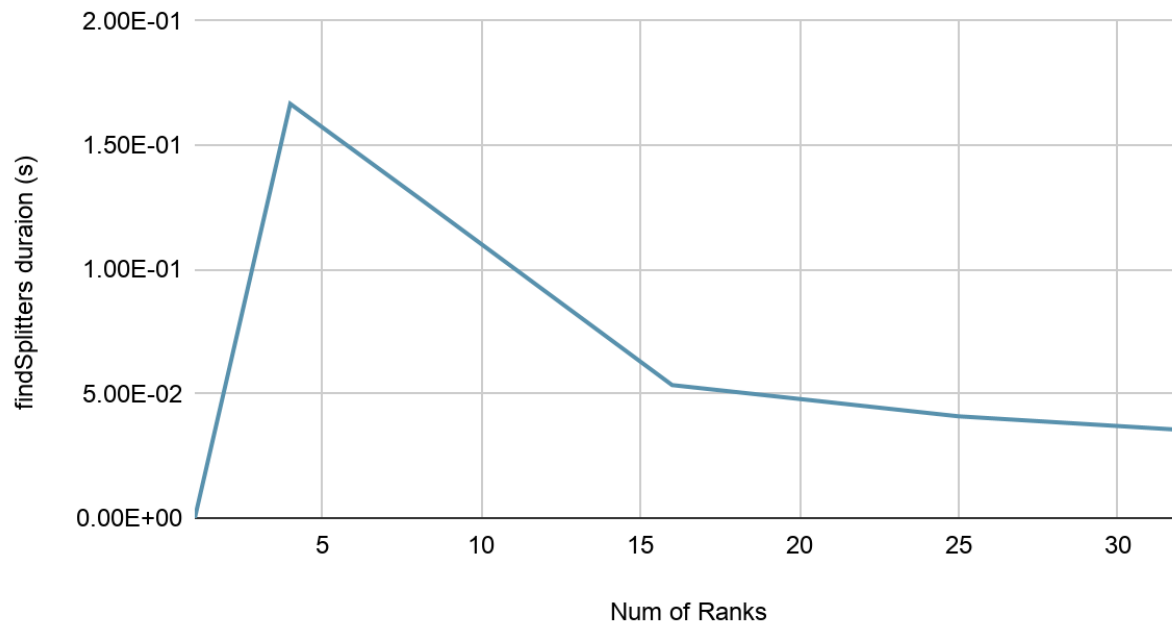


Rebalance Duration vs Num of Ranks



\

findSplitters Duration vs Num of Ranks



findSplitters function ran really fast when only have 1 rank.
This is due to the advantage of not needing to communicate between different ranks, and the input data is already sorted.

moveData Duration vs Num of Ranks

